

# Mounting GDrive directory

```
In [112... from google.colab import drive
drive.mount('/content/gdrive')
%cd '/content/gdrive/MyDrive/Colab Notebooks/UoA_MSDS/Course_8/Capstone1_Ad_
```

Drive already mounted at /content/gdrive; to attempt to forcibly remount, call drive.mount("/content/gdrive", force\_remount=True).  
/content/gdrive/MyDrive/Colab Notebooks/UoA\_MSDS/Course\_8/Capstone1\_Ad\_Campaign\_Recommender

# Import libraries

```
In [113... import os
import csv

from IPython.core.display import display, HTML

import numpy as np
import pandas as pd
from matplotlib import pyplot as plt
import seaborn as sns
%matplotlib inline

import warnings
warnings.filterwarnings('ignore')

import pickle

from sklearn import metrics
from scipy import stats
```

# Reading data

```
In [114... # Load into dataframes, disable quote parsing inside strings when loading
# Ref: https://stackoverflow.com/a/29857126

# Scenario 1: devices with event data
segment_1_train = pd.read_csv('data/train/segment_1_train.csv')
segment_1_test = pd.read_csv('data/test/segment_1_test.csv')

with open('deploy/gender_model_sc1_fit.pickle', 'rb') as f:
    gender_model_sc1_fit = pickle.load(f)
with open('deploy/age_model_sc1_fit.pickle', 'rb') as f:
    age_model_sc1_fit = pickle.load(f)

# Scenario 2: devices without event data
segment_2_train = pd.read_csv('data/train/segment_2_train.csv')
```

```
segment_2_test = pd.read_csv('data/test/segment_2_test.csv')

with open('deploy/gender_model_sc2_fit.pickle', 'rb') as f:
    gender_model_sc2_fit = pickle.load(f)
with open('deploy/age_model_sc2_fit.pickle', 'rb') as f:
    age_model_sc2_fit = pickle.load(f)
```

## Subtask 6: Model evaluation

### 1. Common functions for evaluation

```
In [115... def display_classification_report(y_true, y_pred, clf):
    display(HTML('<h2>Confusion matrix:</h2>'))
    print(metrics.classification_report(
        y_true.values.flatten(),
        y_pred,
        target_names=list(map(lambda cls: f'class {cls}', clf.classes_))),
    )
```

```
In [116... def display_roc_curve(y_true, y_pred):
    fpr, tpr, thresholds = metrics.roc_curve(y_true, y_pred)
    roc_auc = metrics.roc_auc_score(y_true, y_pred)
    display(HTML(f'<h2>AuC score: {roc_auc}</h2>'))
    display(HTML('<h2>ROC curve:</h2>'))
    disp = metrics.RocCurveDisplay(fpr=fpr, tpr=tpr, roc_auc=roc_auc)
    disp.plot()
    plt.show()
```

```
In [117... # Reference: https://towardsdatascience.com/evaluating-classification-models
def display_kolmogorov_smirnov_statistic(y_true, y_pred_proba, clf):
    display(HTML('<h2>Kolmogorov-Smirnov (KS) statistic:</h2>'))
    for cls in clf.classes_:
        df = pd.DataFrame()
        df['real'] = y_true
        df['proba'] = y_pred_proba[:, cls]

        # Recover each class
        target_cls = df[df['real'] == cls]['proba']
        other_cls = df[df['real'] != cls]['proba']
        ks_score = stats.ks_2samp(target_cls, other_cls)
        ks_score = dict(
            statistic=ks_score.statistic,
            pvalue=ks_score.pvalue,
            statistic_location=ks_score.statistic_location,
            statistic_sign=ks_score.statistic_sign
        )
        ks_score = pd.DataFrame(ks_score, index=[0])
        display(HTML(f'<h3>class {cls} vs others:</h3>'))
        display(ks_score)
```

```

In [118... # Reference: https://github.com/vinyluis/Articles/blob/main/Kolmogorov-Smirnov
def display_binary_probability_bands(y_true, y_pred_proba):
    display(HTML('<h2>Probability bands:</h2>'))
    sns.histplot(x=y_pred_proba[:, 1], hue=y_true.values.flatten())

In [119... def display_multi_label_log_loss(y_true, y_pred_proba, clf):
    display(HTML('<h2>Multi-label log loss:</h2>'))
    display(metrics.log_loss(y_true.values.flatten(), y_pred_proba))

In [120... def gender_model_evaluations(X, y, clf):
    y_pred = clf.predict(X)
    y_pred_proba = clf.predict_proba(X)
    display_classification_report(y, y_pred, clf)
    display_roc_curve(y, y_pred)
    display_kolmogorov_smirnov_statistic(y, y_pred_proba, clf)
    display_binary_probability_bands(y, y_pred_proba)

In [121... def age_model_evaluations(X, y, clf):
    y_pred = clf.predict(X)
    y_pred_proba = clf.predict_proba(X)
    display_classification_report(y, y_pred, clf)
    display_kolmogorov_smirnov_statistic(y, y_pred_proba, clf)
    display_multi_label_log_loss(y, y_pred_proba, clf)

```

## 2. Gender prediction model

### a) Scenario 1: Devices with event data

```

In [122... # Format train/test data
X_train = segment_1_train.drop(columns=['device_id', 'gender', 'age_group'])
y_train = segment_1_train[['gender']]
X_test = segment_1_test.drop(columns=['device_id', 'gender', 'age_group'])
y_test = segment_1_test[['gender']]

for df, label in zip(
    [
        X_train,
        y_train,
        X_test,
        y_test,
    ],
    [
        'X_train',
        'y_train',
        'X_test',
        'y_test',
    ],
):
    display(HTML(f'<h2>{label}</h2>'))
    display(df)

```

# X\_train

	average_daily_events	location_cluster_-1	location_cluster_0	location_cluster_1	
0	0.380427	1.0	0.0	0.0	
1	-0.621224	1.0	0.0	0.0	
2	-2.669088	1.0	0.0	0.0	
3	1.157734	1.0	0.0	0.0	
4	-0.123856	1.0	0.0	0.0	
...	...	...	...	...	
2836	-0.623716	1.0	0.0	0.0	
2837	1.288079	0.0	0.0	0.0	
2838	-0.705781	1.0	0.0	0.0	
2839	-0.026953	1.0	0.0	0.0	
2840	1.041286	0.0	1.0	0.0	

2841 rows x 1856 columns

# y\_train

	gender
0	1
1	1
2	1
3	1
4	0
...	...
2836	1
2837	1
2838	0
2839	1
2840	1

2841 rows x 1 columns

# X\_test

	average_daily_events	location_cluster_-1	location_cluster_0	location_cluster_1	location_cluster_2
0	2.356397	0.0	0.0	0.0	0.0
1	-0.523348	1.0	0.0	0.0	0.0
2	0.073159	1.0	0.0	0.0	0.0
3	1.744514	0.0	0.0	0.0	0.0
4	1.057478	1.0	0.0	0.0	0.0
...	...	...	...	...	...
1213	0.480375	1.0	0.0	0.0	0.0
1214	-0.893056	1.0	0.0	0.0	0.0
1215	0.478324	1.0	0.0	0.0	0.0
1216	-0.013296	0.0	0.0	0.0	0.0
1217	0.272121	1.0	0.0	0.0	0.0

1218 rows × 1856 columns

# y\_test

	gender
0	0
1	1
2	1
3	1
4	0
...	...
1213	0
1214	0
1215	1
1216	1
1217	0

1218 rows × 1 columns

## Evaluations on train set:


```
In [123... gender_model_evaluations(  
    X_train,  
    y_train,  
    gender_model_sc1_fit,  
)
```

## Confusion matrix:

	precision	recall	f1-score	support
class 0	1.00	0.01	0.02	791
class 1	0.72	1.00	0.84	2050
accuracy			0.72	2841
macro avg	0.86	0.51	0.43	2841
weighted avg	0.80	0.72	0.61	2841

AuC score: 0.5050568900126422

## ROC curve:



## Kolmogorov–Smirnov (KS) statistic:


class 0 vs others:

	statistic	pvalue	statistic_location	statistic_sign
0	0.312131	1.087792e-49	0.274235	-1

class 1 vs others:

	statistic	pvalue	statistic_location	statistic_sign
0	0.312131	1.087792e-49	0.725739	-1

Probability bands:



Evaluations on test set:


```
In [124... gender_model_evaluations(  
    X_test,  
    y_test,  
    gender_model_sc1_fit,  
)
```

Confusion matrix:

	precision	recall	f1-score	support
class 0	0.67	0.01	0.01	303
class 1	0.75	1.00	0.86	915
accuracy			0.75	1218
macro avg	0.71	0.50	0.44	1218
weighted avg	0.73	0.75	0.65	1218

**AuC score: 0.5027538819455716**

**ROC curve:**



**Kolmogorov–Smirnov (KS) statistic:**

**class 0 vs others:**


	statistic	pvalue	statistic_location	statistic_sign
0	0.189666	1.207037e-07	0.285835	-1

**class 1 vs others:**



	statistic	pvalue	statistic_location	statistic_sign
0	0.189666	1.207037e-07	0.714125	-1

## Probability bands:



## b) Scenario 2: Devices without event data

```
In [125... # Format train/test data
X_train = segment_2_train.drop(columns=['device_id', 'gender', 'age_group'])
y_train = segment_2_train[['gender']]
X_test = segment_2_test.drop(columns=['device_id', 'gender', 'age_group'])
y_test = segment_2_test[['gender']]

for df, label in zip(
    [
        X_train,
        y_train,
        X_test,
        y_test,
    ],
    [
        'X_train',
        'y_train',
        'X_test',
        'y_test',
    ],
)
```

```
display(HTML(f'<h2>{label}</h2>'))
display(df)
```

## X\_train

	phone_brand_AUX	phone_brand_Bacardi	phone_brand_Bifer	phone_brand_CUI
0	0.0	0.0	0.0	0
1	0.0	0.0	0.0	0
2	0.0	0.0	0.0	0
3	0.0	0.0	0.0	0
4	0.0	0.0	0.0	0
...	...	...	...	
49541	0.0	0.0	0.0	0
49542	0.0	0.0	0.0	0
49543	0.0	0.0	0.0	0
49544	0.0	0.0	0.0	0
49545	0.0	0.0	0.0	0

49546 rows × 97 columns

## y\_train

	gender
0	1
1	1
2	1
3	1
4	1
...	...
49541	0
49542	1
49543	1
49544	1
49545	0

49546 rows × 1 columns

# X\_test

	phone_brand_AUX	phone_brand_Bacardi	phone_brand_Bifer	phone_brand_CUB
0	0.0	0.0	0.0	0.
1	0.0	0.0	0.0	0.
2	0.0	0.0	0.0	0.
3	0.0	0.0	0.0	0.
4	0.0	0.0	0.0	0.
...	...	...	...	...
21230	0.0	0.0	0.0	0.
21231	0.0	0.0	0.0	0.
21232	0.0	0.0	0.0	0.
21233	0.0	0.0	0.0	0.
21234	0.0	0.0	0.0	0.

21235 rows × 97 columns

# y\_test

	gender
0	0
1	1
2	0
3	0
4	1
...	...
21230	1
21231	0
21232	0
21233	1
21234	1

21235 rows × 1 columns

Evaluations on train set:


```
In [126... gender_model_evaluations(
    X_train,
    y_train,
    gender_model_sc2_fit,
)
```

## Confusion matrix:

	precision	recall	f1-score	support
class 0	0.71	0.01	0.02	18035
class 1	0.64	1.00	0.78	31511
accuracy			0.64	49546
macro avg	0.67	0.50	0.40	49546
weighted avg	0.66	0.64	0.50	49546

AuC score: 0.5029333044237649

## ROC curve:



## Kolmogorov–Smirnov (KS) statistic:


class 0 vs others:

	statistic	pvalue	statistic_location	statistic_sign
0	0.101814	5.762835e-104	0.347641	-1

class 1 vs others:

	statistic	pvalue	statistic_location	statistic_sign
0	0.101814	5.762835e-104	0.651273	-1

Probability bands:



Evaluations on test set:


```
In [127... gender_model_evaluations(
    X_test,
    y_test,
    gender_model_sc2_fit,
)
```

Confusion matrix:

	precision	recall	f1-score	support
class 0	0.61	0.01	0.01	7676
class 1	0.64	1.00	0.78	13559
accuracy			0.64	21235
macro avg	0.62	0.50	0.40	21235
weighted avg	0.63	0.64	0.50	21235

**AuC score: 0.5021051388291212**

**ROC curve:**



**Kolmogorov–Smirnov (KS) statistic:**


**class 0 vs others:**

	statistic	pvalue	statistic_location	statistic_sign
0	0.105391	7.443207e-48	0.347641	-1

**class 1 vs others:**

	statistic	pvalue	statistic_location	statistic_sign
0	0.105391	7.443207e-48	0.651273	-1

## Probability bands:



## 3. Age group prediction model

### a) Scenario 1: Devices with event data

```
In [128... # Format train/test data
X_train = segment_1_train.drop(columns=['device_id', 'gender', 'age_group'])
y_train = segment_1_train[['age_group']]
X_test = segment_1_test.drop(columns=['device_id', 'gender', 'age_group'])
y_test = segment_1_test[['age_group']]

for df, label in zip(
    [
        X_train,
        y_train,
        X_test,
        y_test,
    ],
    [
        'X_train',
        'y_train',

```

```
        'y_test',
    ],
):
    display(HTML(f'<h2>{label}</h2>'))
    display(df)
```

## X\_train

	average_daily_events	location_cluster_-1	location_cluster_0	location_cluster_1	
0	0.380427	1.0	0.0	0.0	
1	-0.621224	1.0	0.0	0.0	
2	-2.669088	1.0	0.0	0.0	
3	1.157734	1.0	0.0	0.0	
4	-0.123856	1.0	0.0	0.0	
...	...	...	...	...	
2836	-0.623716	1.0	0.0	0.0	
2837	1.288079	0.0	0.0	0.0	
2838	-0.705781	1.0	0.0	0.0	
2839	-0.026953	1.0	0.0	0.0	
2840	1.041286	0.0	1.0	0.0	

2841 rows × 1856 columns

## y\_train



age_group	
0	2
1	2
2	0
3	2
4	1
...	...
2836	2
2837	1
2838	0
2839	2
2840	2

2841 rows × 1 columns

## X\_test

	average_daily_events	location_cluster_-1	location_cluster_0	location_cluster_1	location_cluster_2
0	2.356397	0.0	0.0	0.0	0.0
1	-0.523348	1.0	0.0	0.0	0.0
2	0.073159	1.0	0.0	0.0	0.0
3	1.744514	0.0	0.0	0.0	0.0
4	1.057478	1.0	0.0	0.0	0.0
...	...	...	...	...	...
1213	0.480375	1.0	0.0	0.0	0.0
1214	-0.893056	1.0	0.0	0.0	0.0
1215	0.478324	1.0	0.0	0.0	0.0
1216	-0.013296	0.0	0.0	0.0	0.0
1217	0.272121	1.0	0.0	0.0	0.0

1218 rows × 1856 columns

## y\_test

age_group	
0	3
1	1
2	1
3	3
4	2
...	...
1213	1
1214	2
1215	1
1216	1
1217	1

1218 rows × 1 columns

Evaluations on train set:

```
In [129... age_model_evaluations(  
    X_train,  
    y_train,  
    age_model_sc1_fit,  
)
```

Confusion matrix:

	precision	recall	f1-score	support
class 0	0.00	0.00	0.00	455
class 1	0.45	0.84	0.59	1086
class 2	0.47	0.41	0.44	922
class 3	0.00	0.00	0.00	378
accuracy			0.46	2841
macro avg	0.23	0.31	0.26	2841
weighted avg	0.33	0.46	0.37	2841

Kolmogorov–Smirnov (KS) statistic:

class 0 vs others:

	statistic	pvalue	statistic_location	statistic_sign
0	0.402283	4.717071e-56	0.208638	-1

### class 1 vs others:

	statistic	pvalue	statistic_location	statistic_sign
0	0.298434	3.148953e-53	0.3245	-1

### class 2 vs others:

	statistic	pvalue	statistic_location	statistic_sign
0	0.272213	3.986403e-41	0.291119	-1

### class 3 vs others:

	statistic	pvalue	statistic_location	statistic_sign
0	0.427447	1.373441e-54	0.205205	-1

## Multi-label log loss:

1.266620233023434

### Evaluations on test set:

```
In [130... age_model_evaluations(
    X_test,
    y_test,
    age_model_sc1_fit,
)
```

## Confusion matrix:

	precision	recall	f1-score	support
class 0	0.00	0.00	0.00	190
class 1	0.39	0.78	0.52	449
class 2	0.43	0.34	0.38	423
class 3	0.00	0.00	0.00	156
accuracy			0.40	1218
macro avg	0.21	0.28	0.23	1218
weighted avg	0.29	0.40	0.32	1218

## Kolmogorov–Smirnov (KS) statistic:

### class 0 vs others:

	statistic	pvalue	statistic_location	statistic_sign
0	0.25084	2.269936e-09	0.180566	-1

### class 1 vs others:

	statistic	pvalue	statistic_location	statistic_sign
0	0.090196	0.018347	0.31169	-1

### class 2 vs others:

	statistic	pvalue	statistic_location	statistic_sign
0	0.1499	0.000007	0.287722	-1

### class 3 vs others:

	statistic	pvalue	statistic_location	statistic_sign
0	0.188867	0.000101	0.195758	-1

## Multi-label log loss:

1.2957567985120026

### b) Scenario 2: Devices without event data

```
In [131]: # Format train/test data
X_train = segment_2_train.drop(columns=['device_id', 'gender', 'age_group'])
y_train = segment_2_train[['age_group']]
X_test = segment_2_test.drop(columns=['device_id', 'gender', 'age_group'])
y_test = segment_2_test[['age_group']]

for df, label in zip(
    [
        X_train,
        y_train,
        X_test,
        y_test,
    ],
    [
        'X_train',
        'y_train',
        'X_test',
        'y_test',
    ],
):
```

```
display(HTML(f'<h2>{label}</h2>'))
display(df)
```

X\_train

	phone_brand_AUX	phone_brand_Bacardi	phone_brand_Bifer	phone_brand_CUI
0	0.0	0.0	0.0	0
1	0.0	0.0	0.0	0
2	0.0	0.0	0.0	0
3	0.0	0.0	0.0	0
4	0.0	0.0	0.0	0
...	...	...	...	
49541	0.0	0.0	0.0	0
49542	0.0	0.0	0.0	0
49543	0.0	0.0	0.0	0
49544	0.0	0.0	0.0	0
49545	0.0	0.0	0.0	0

49546 rows × 97 columns

y\_train

	age_group
0	1
1	0
2	1
3	1
4	1
...	...
49541	2
49542	0
49543	0
49544	1
49545	2

49546 rows × 1 columns

# X\_test

	phone_brand_AUX	phone_brand_Bacardi	phone_brand_Bifer	phone_brand_CUB
0	0.0	0.0	0.0	0.
1	0.0	0.0	0.0	0.
2	0.0	0.0	0.0	0.
3	0.0	0.0	0.0	0.
4	0.0	0.0	0.0	0.
...	...	...	...	.
21230	0.0	0.0	0.0	0.
21231	0.0	0.0	0.0	0.
21232	0.0	0.0	0.0	0.
21233	0.0	0.0	0.0	0.
21234	0.0	0.0	0.0	0.

21235 rows × 97 columns

# y\_test

	age_group
0	3
1	2
2	1
3	2
4	2
...	...
21230	1
21231	1
21232	0
21233	0
21234	0

21235 rows × 1 columns

Evaluations on train set:

```
In [132]: age_model_evaluations(  
        X_train,  
        y_train,  
        age_model_sc2_fit,  
    )
```

Confusion matrix:

	precision	recall	f1-score	support
class 0	0.00	0.00	0.00	12185
class 1	0.41	1.00	0.59	20488
class 2	0.00	0.00	0.00	12249
class 3	0.00	0.00	0.00	4624
accuracy			0.41	49546
macro avg	0.10	0.25	0.15	49546
weighted avg	0.17	0.41	0.24	49546

Kolmogorov–Smirnov (KS) statistic:

class 0 vs others:

	statistic	pvalue	statistic_location	statistic_sign
0	0.132909	5.394860e-142	0.248211	-1

class 1 vs others:

	statistic	pvalue	statistic_location	statistic_sign
0	0.040295	2.172174e-17	0.306779	-1

class 2 vs others:

	statistic	pvalue	statistic_location	statistic_sign
0	0.12387	8.984869e-124	0.248906	-1

class 3 vs others:

	statistic	pvalue	statistic_location	statistic_sign
0	0.117302	1.034283e-50	0.202477	-1

Multi-label log loss:

1.321135864631613

## Evaluations on test set:

```
In [133... age_model_evaluations(  
    X_test,  
    y_test,  
    age_model_sc2_fit,  
)
```

## Confusion matrix:

	precision	recall	f1-score	support
class 0	0.00	0.00	0.00	5272
class 1	0.41	1.00	0.58	8619
class 2	0.00	0.00	0.00	5356
class 3	0.00	0.00	0.00	1988
accuracy			0.41	21235
macro avg	0.10	0.25	0.14	21235
weighted avg	0.16	0.41	0.23	21235

## Kolmogorov–Smirnov (KS) statistic:

### class 0 vs others:

	statistic	pvalue	statistic_location	statistic_sign
0	0.134013	1.570777e-62	0.248211	-1

### class 1 vs others:

	statistic	pvalue	statistic_location	statistic_sign
0	0.04522	1.543816e-09	0.301402	-1

### class 2 vs others:

	statistic	pvalue	statistic_location	statistic_sign
0	0.120592	3.275415e-51	0.241536	-1

### class 3 vs others:

	statistic	pvalue	statistic_location	statistic_sign
0	0.129276	1.024451e-26	0.199844	-1

## Multi-label loss:



## 4. Conclusion

All the models perform badly and fail to effectively separate the prediction classes although the accuracy seems to be good. The classification reports show that there is not much difference between the scores of train and test set, so the models are not overfit.

Looking at ROC score for gender prediction, it's just slightly better than random guess. This is further proved by KS test, which shows low statistic for every class.

For age group prediction, the log loss is also high, this shows that the predictions from models is far from actual labels.

For gender and age group predictions, models for scenario 1 (where we have event data available) show better performance than models for scenario 2. For deployment, we better use models for scenario 1 only.

Due to the quick tuning during model training phase, the final models are not yet the best ones, so if we want better prediction, we need to tune using grid search to try all the combinations of hyper parameters. The wait time is very long if using Google Colab, so we will just skip for now and use the models at hand for deployment.

For picking the threshold or deciles for predictions, as there is no clear separation gap between the classes (see the plot of probability bands), we can't have any sign for a good threshold or value of deciles to use. For deployment, we will just stick to the class with highest probability.